

A*pex: Efficient Approximate Multi-Objective Search on Graphs



Han Zhang (zhan645@usc.edu),¹ Oren Salzman,² T. K. Satish Kumar,¹ Ariel Felner,³ Carlos Hernández Ulloa,⁴ Sven Koenig¹

¹ University of Southern California, ² Technion - Israel Institute of Technology, ³ Ben-Gurion University, ⁴ Universidad San Sebastian

1 Background




Figure: An example search instance. Blue and red arrows show the two paths in the Pareto-optimal frontier.

Definition (Multi-objective search).

A *multi-objective search instance* contains a graph, a start state, and a goal state. A *cost function* c maps an edge in the graph to a cost vector of length N . A *solution path* π is a path from the start state to the goal state. The *cost of path* π , denoted as $c(\pi)$, is the accumulated costs of its edges.

Definition (Weak dominance and ϵ -dominance).

Path π *weakly dominates* path π' iff $c(\pi) \leq c(\pi')$. For a non-negative real number ϵ , path π ϵ -*dominates* path π' iff $c(\pi) \leq (1 + \epsilon)c(\pi')$.


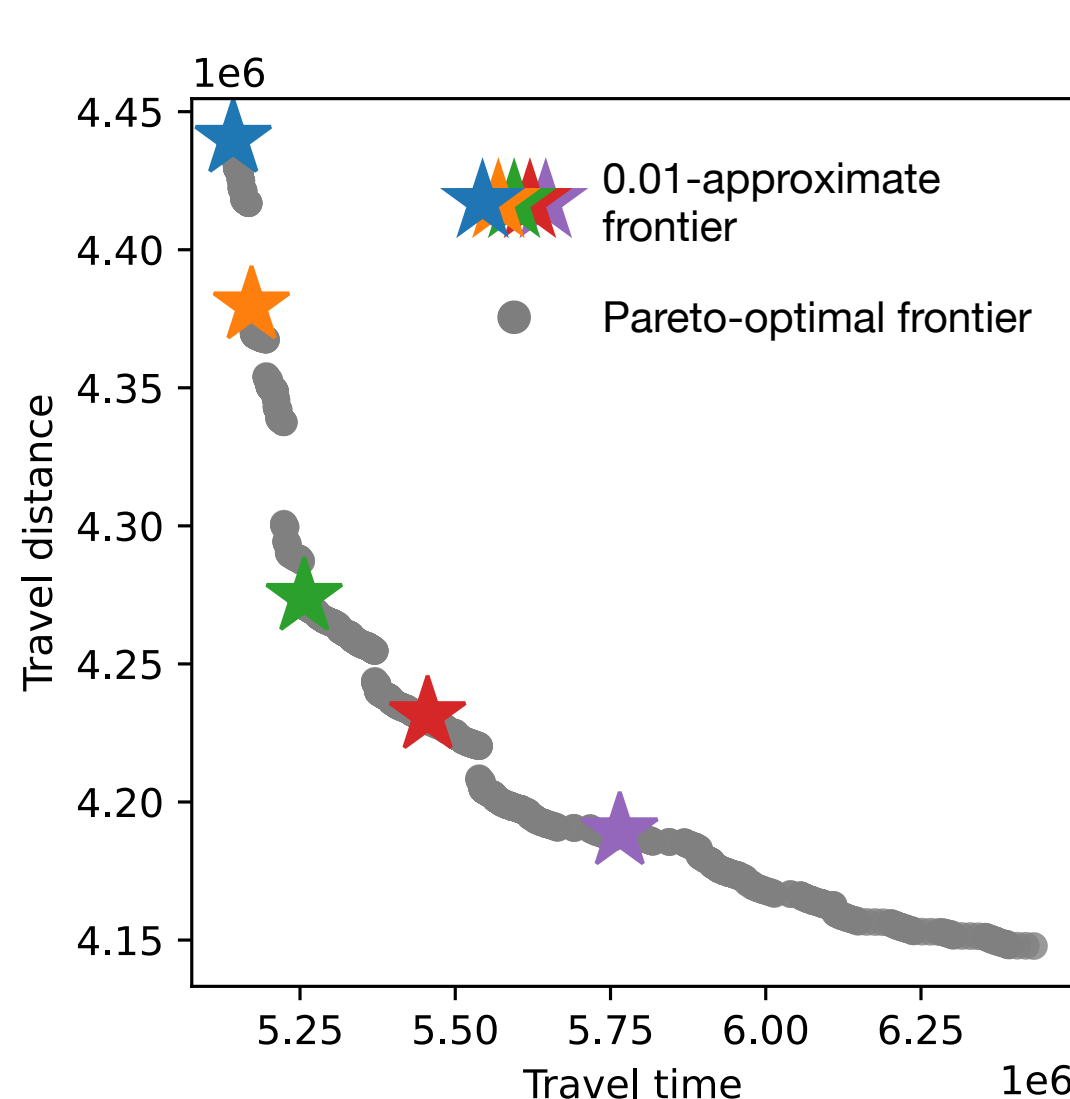
Definition (Pareto-optimal frontier).

A set of solution paths Π is a *Pareto-optimal frontier* iff (1) every solution path for the problem instance is weakly dominated by at least one solution path in Π and (2) solution paths in Π do not weakly dominate each other.

Definition (Approximate Pareto-optimal frontier)

Given an ϵ -value, a set of solution paths Π_ϵ is an ϵ -*approximate frontier* iff (1) every solution path for the problem instance is ϵ -dominated by at least one solution path in Π_ϵ and (2) solution paths in Π_ϵ do not weakly dominate each other.

4 ϵ -Approximate Frontier: An Example



The left figure shows the Pareto-optimal frontier and an ϵ ($=0.01$)-approximate frontier for a road-network problem instance with two objectives. The right figure shows the corresponding paths of the ϵ -approximate frontier.

In this example, the Pareto-optimal frontier has **many** ($=2253$) **solution paths** with many of them being very similar. The ϵ -approximate frontier contains only five diverse solution paths and takes a **much smaller amount of time** ($\sim 5\%$) to compute.

5 A*pex

In A*pex, each search node is an *apex-path pair*, which contains a *representative path* π and a cost vector A called *apex*. An apex-path pair is ϵ -*bounded* iff $c(\pi) \leq (1 + \epsilon)A$.






Figure: An apex-path pair

A*pex runs an A*-like search with apex-path pairs as search nodes and merges apex-path pairs whenever the resulted apex-path pairs are ϵ -bounded. The solution set returned by A*pex contains the representative path of each solution apex-path pair.

Summary: By using ϵ -bounded apex-path pairs to approximate the Pareto-optimal frontier, A*pex finds an ϵ -approximate frontier with fewer node expansions than existing algorithms.



(a) Extending an apex-path pair



(b) Merging two apex-path pairs

Figure: Operations on apex-path pairs

2 Motivation

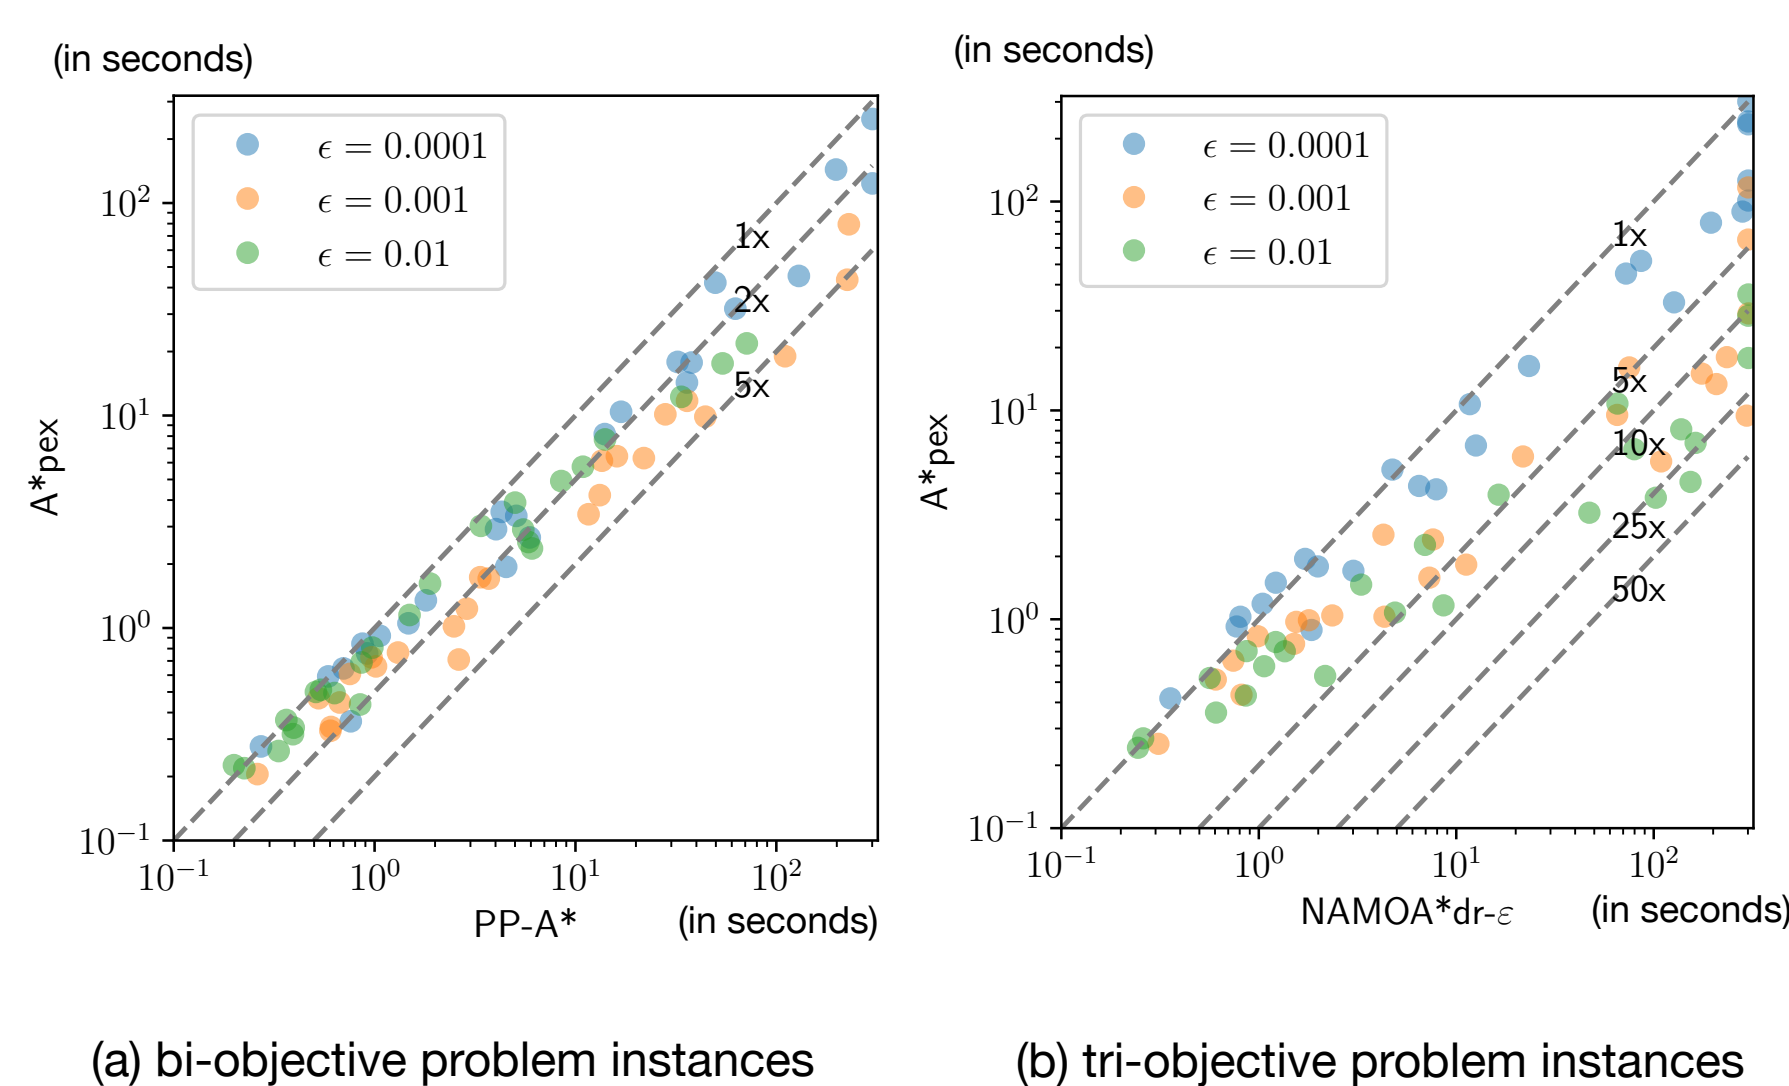
- ▶ The Pareto-optimal frontier can be **exponential** in the size of the graph being searched, which makes multi-objective search time-consuming.
- ▶ The ϵ -approximate frontiers are typically much smaller than the Pareto-optimal frontiers even for small ϵ -values and hence easier to compute.
- ▶ The existing algorithm **PP-A*** [1] finds ϵ -approximate frontiers for bi-objective search instances efficiently. However, it is unclear how to generalize PP-A* to search instances with more than two objectives.

3 Contributions

We proposed an approximate multi-objective search algorithm **A*pex**, which:

- ▶ finds an ϵ -approximate frontier for a user-provided ϵ -value,
- ▶ builds on PP-A* and also makes it **more efficient** for bi-objective search, and
- ▶ can solve problem instances with **more than two** objectives while PP-A* cannot.

6 Experimental Results



The figures show the runtime of A*pex and baseline algorithms for different numbers of objectives, problem instances, and ϵ -values. The **y-coordinate** of each point is the runtime of A*pex, and the **x-coordinate** of each point is the runtime of the respective baseline algorithm.

Summary: A*pex runs faster than the respective baseline algorithm in most problem instances for both two and three objectives. In some instances, the speedups are more than **5x** and **25x**, respectively.