


# BEAM SEARCH: FASTER & MONOTONIC

Sofia Lemons,<sup>1, 2</sup> Carlos Linares López,<sup>3</sup> Robert C. Holte,<sup>4</sup> Wheeler Ruml<sup>1</sup>

<sup>1</sup> University of New Hampshire, <sup>2</sup> Earlham College, <sup>3</sup> Universidad Carlos III de Madrid,  
<sup>4</sup> University of Alberta, Alberta Machine Intelligence Institute (Amii)

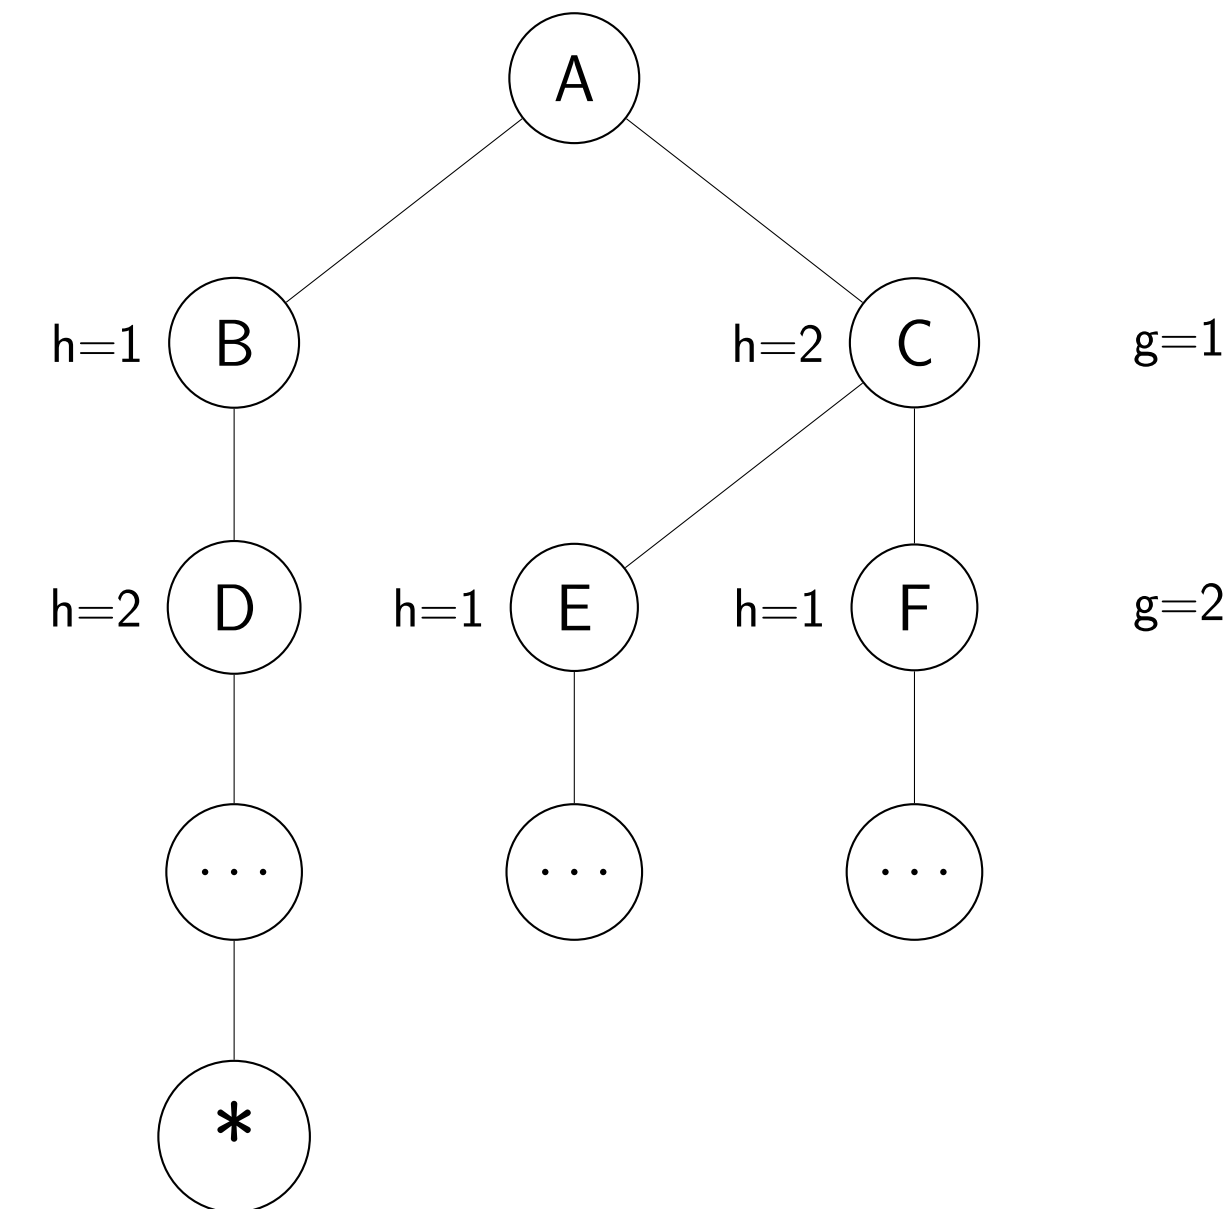
## Non-monotonicity



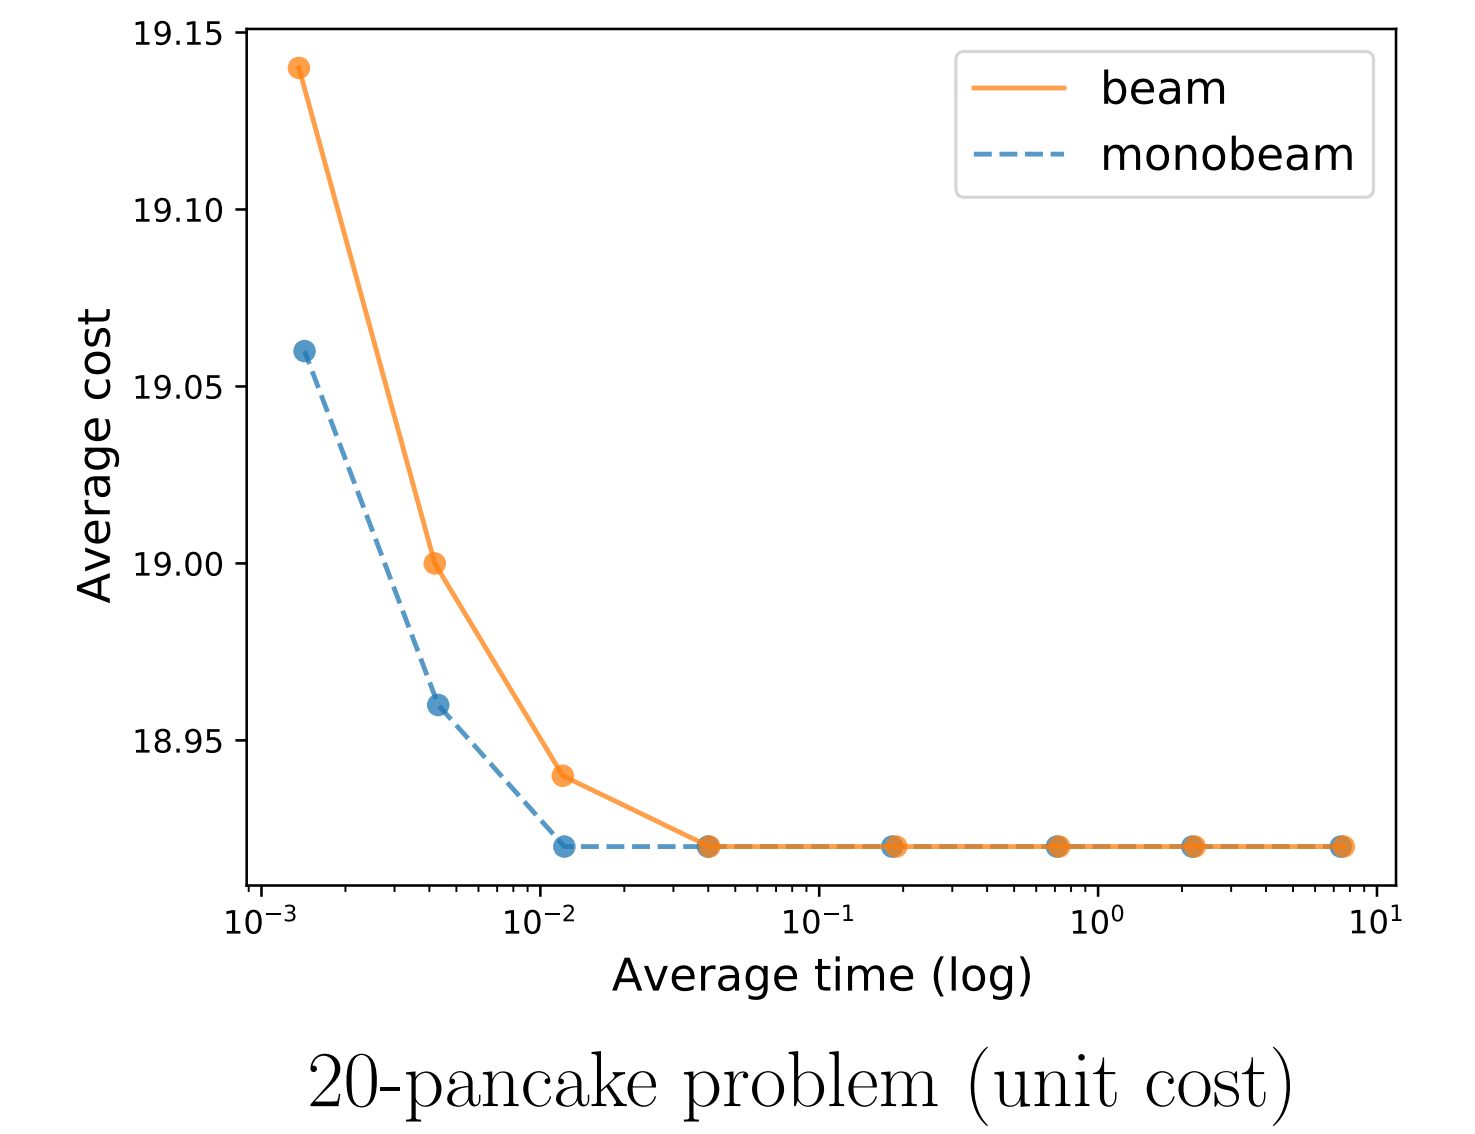
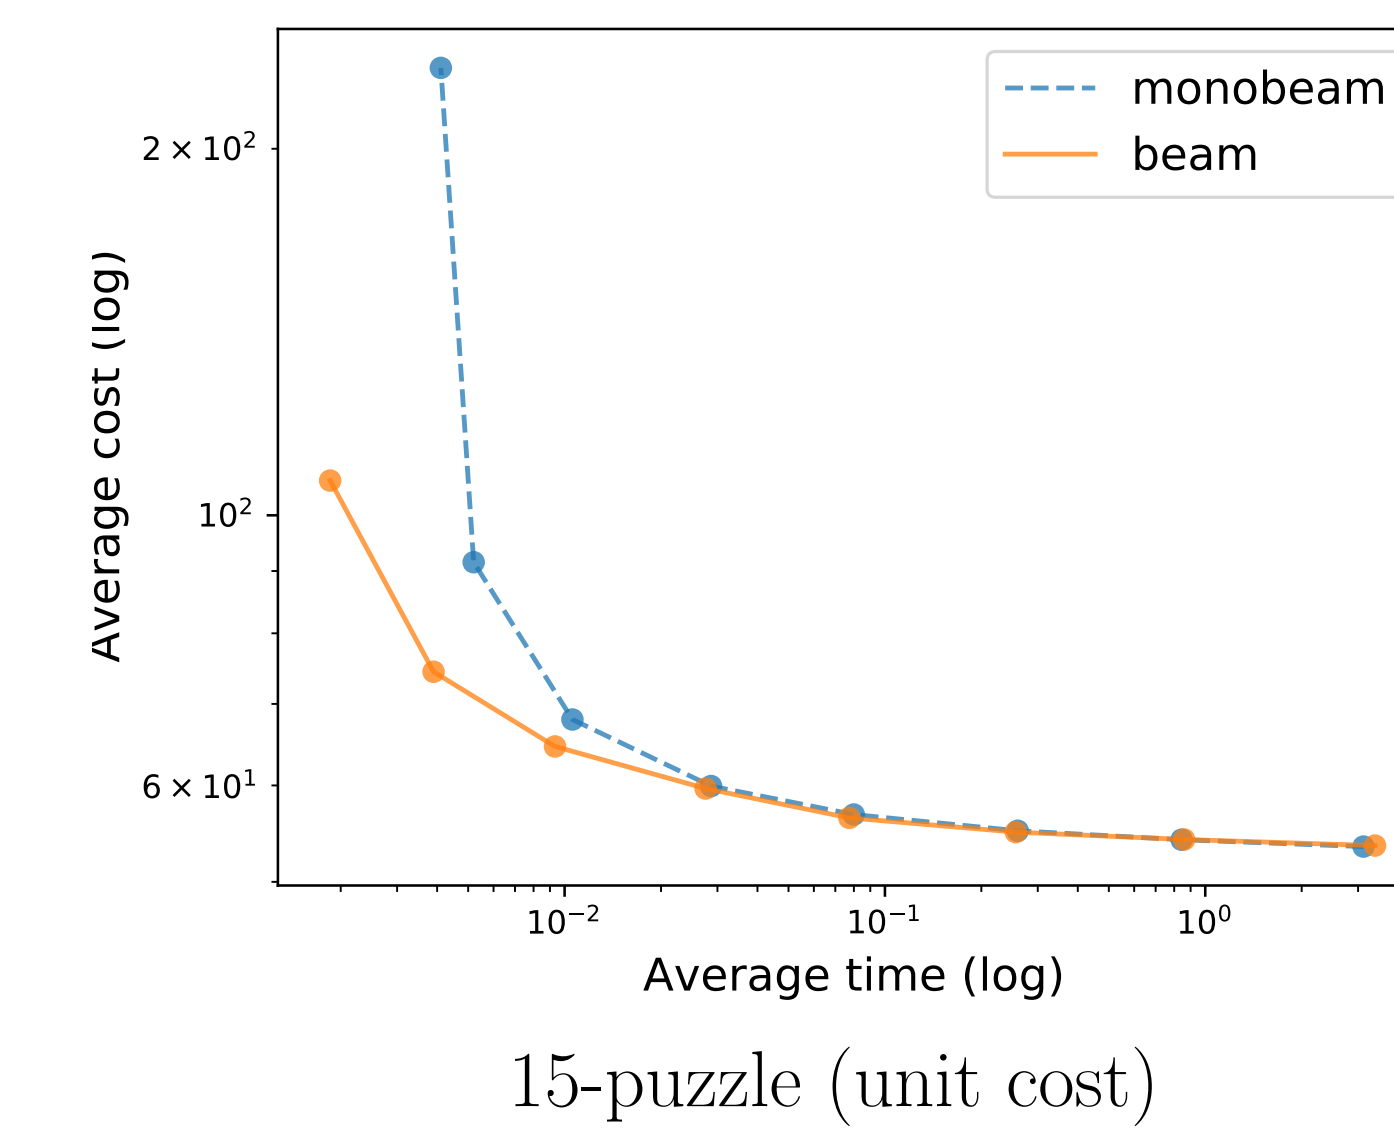
When increasing beam width, beam search sometimes returns worse solutions.

## Root Cause: Cuckoo Nodes

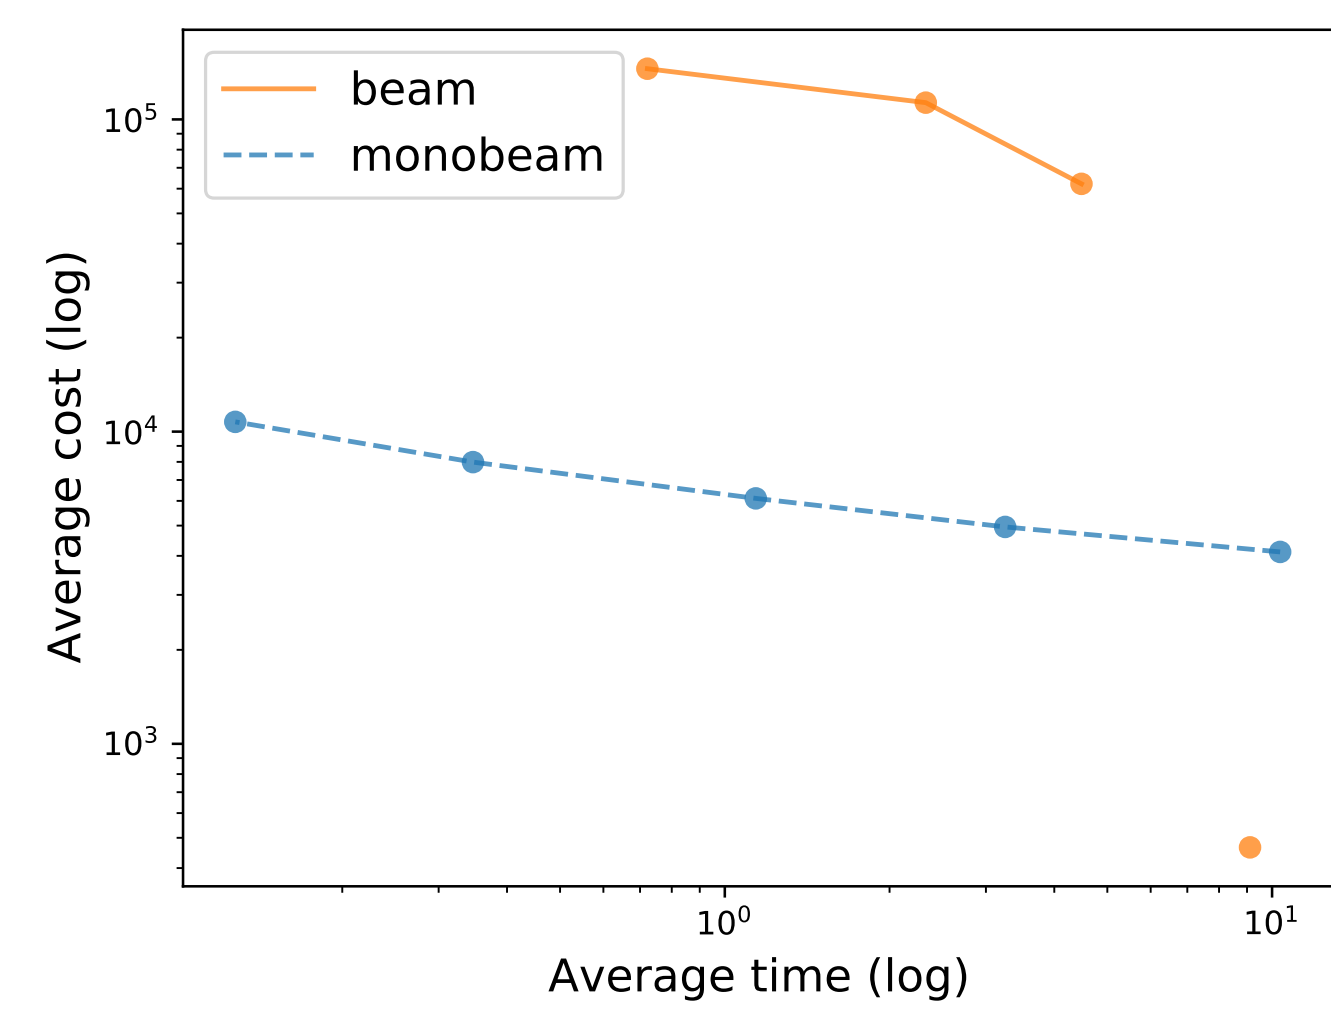
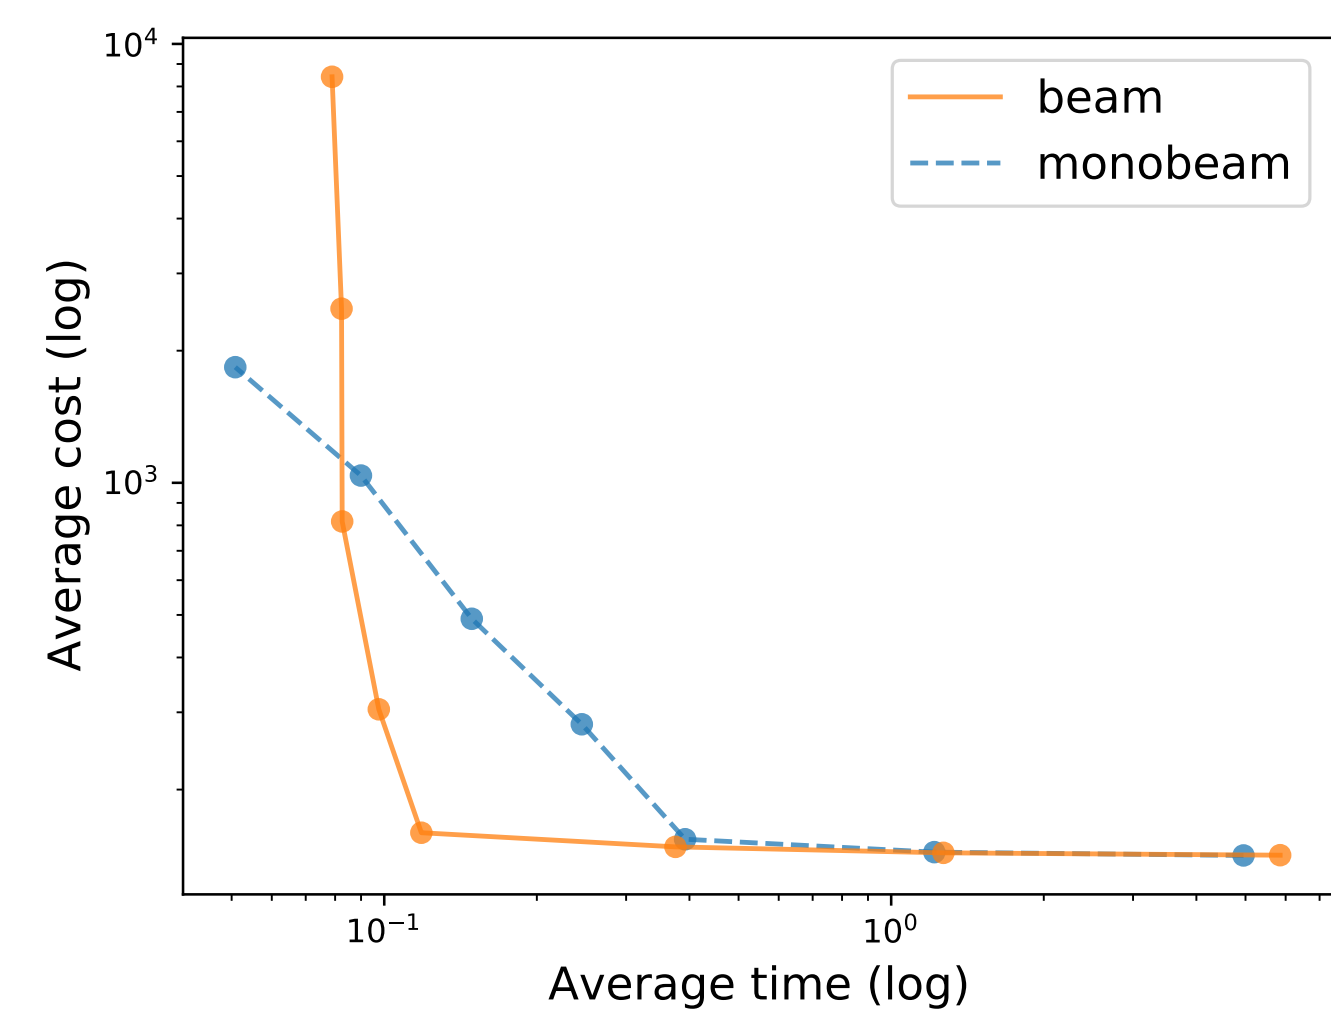
The children of nodes generated at the larger beam width may have incorrectly low cost-to-go estimates due to heuristic error. We refer to these as *cuckoo nodes* because they cause truly better nodes to be pushed off the beam, just as, when a cuckoo bird lays its eggs in another bird's nest, the cuckoo chicks push the other bird's eggs out of the nest. In the example, a beam search with width 2 returns a solution of lower quality than a beam search with width 1. When the width of the beam is increased, node C serves as a cuckoo node, with its successors pushing the successors of B off the beam.



## Results (Monobeam)




## Trouble in Non-Unit Cost



In non-unit cost, beam and monobeam often find poor solutions or fail to solve.

## Monobeam

In order to avoid cuckoo nodes as the beam is widened, **monobeam** considers each specific position in the beam, or *slot*, sequentially. When selecting a node to fill slot  $s$  on the beam, considers only successors of the nodes currently in slots 1 through  $s$ . This preserves the search order in earlier slots, regardless of increases in beam width.



Monobeam with width  $w$  returns equal or better solution than any width lower than  $w$ .

## Searching on Distance-to-go

For best-first search, searching using distance-to-go  $d$  instead of cost-to-go  $h$  is known to yield faster search in non-unit cost domains. It is natural to ask whether guiding beam search using  $d$  might help. We introduce **bead** search, which is beam search using a purely distance-based measurement  $l(n) = \text{depth}(n) + d(n)$  (estimated length of solution.) We also introduce **monobead** search, which is monotonic beam search, except using  $l(n)$  to select nodes for the beam.

## Conclusions

Monobeam guarantees solution cost monotonic in beam width.

The guarantee of monotonicity sometimes comes at the expense of slightly worse average solution cost, but sometimes better.

Using distance-to-go instead of cost-to-go provided significant improvement in beam searches.

## Results (Bead & Monobead)

