


# Reinforcement Learning for Classical Planning: Viewing Heuristics as Dense Reward Generators

Clement Gehring\*, Masataro Asai\*, Rohan Chitnis, Tom Silver, Leslie Pack Kaelbling, Shirin Sohrabi, Michael Katz

Corresponding Author: clement@gehring.io



## Abstract

Recent advances in reinforcement learning (RL) have led to a growing interest in applying RL to classical planning domains or applying classical planning methods to some complex RL domains. However, the long-horizon goal-based problems found in classical planning lead to sparse rewards for RL, making direct application inefficient. In this paper, we propose to leverage domain-independent heuristic functions commonly used in the classical planning literature to improve the sample efficiency of RL. These classical heuristics act as dense reward generators to alleviate the sparse-rewards issue and enable our RL agent to learn domain-specific value functions as residuals on these heuristics, making learning easier. Correct application of this technique requires consolidating the discounted metric used in RL and the non-discounted metric used in heuristics. We implement the value functions using Neural Logic Machines, a neural network architecture designed for grounded first-order logic inputs. We demonstrate on several classical planning domains that using classical heuristics for RL allows for good sample efficiency compared to sparse-reward RL. We further show that our learned value functions generalize to novel problem instances in the same domain. The source code is available at [github.com/ibm/pddlrl](https://github.com/ibm/pddlrl)

## Coverage Results

| domain (total)  | Baselines          |                  |                 | Ours (mean (max) of 20 runs) |                        |                  | GBFS     |        |
|-----------------|--------------------|------------------|-----------------|------------------------------|------------------------|------------------|----------|--------|
|                 | $h^{\text{blind}}$ | $h^{\text{add}}$ | $h^{\text{FF}}$ | $H^{\text{blind}}$           | $H^{\text{add}}$       | $H^{\text{FF}}$  | -HGN [4] | -H [3] |
| blocks (250)    | 0                  | 126              | 87              | <b>73.1 (94)</b>             | <b>186.6 (229)</b>     | <b>104 (114)</b> | 3        | 208    |
| ferry (250)     | 0                  | 138              | 250             | <b>40.4 (62)</b>             | <b>233.9 (249)</b>     | 250 (250)        | 27       | 240    |
| gripper (250)   | 0                  | 250              | 250             | <b>47.5 (85)</b>             | 250 (250)              | 250 (250)        | 63       | 139    |
| logistics (250) | 0                  | 106              | 243             | 0 (0)                        | 54.1 6.8( <b>115</b> ) | 79.8 12.9(189)   | -        | 0      |
| miconic (442)   | 171                | 442              | 442             | 143.3 ( <b>246</b> )         | 442 (442)              | 440.8 (442)      | -        | 0      |
| parking (700)   | 0                  | 607              | 700             | <b>0.9 (3)</b>               | <b>619 (689)</b>       | 696.9 (700)      | -        | 333    |
| satellite (250) | 0                  | 249              | 222             | <b>26.5 (99)</b>             | 233.3 (250)            | 163.2 (205)      | -        | 9      |
| visittall (252) | 252                | 252              | 252             | 207.6 (238)                  | 251.9 (252)            | 252 (252)        | -        | 101    |

Table 1: A comparison of the number of test instances solved. **Note:** the method labelled GBFS-HGN [4] was designed in the context of  $A^*$  but was used with GBFS in our work. Therefore, these results are not indicative of the capabilities of this method when using other search algorithms.

## Our Goal

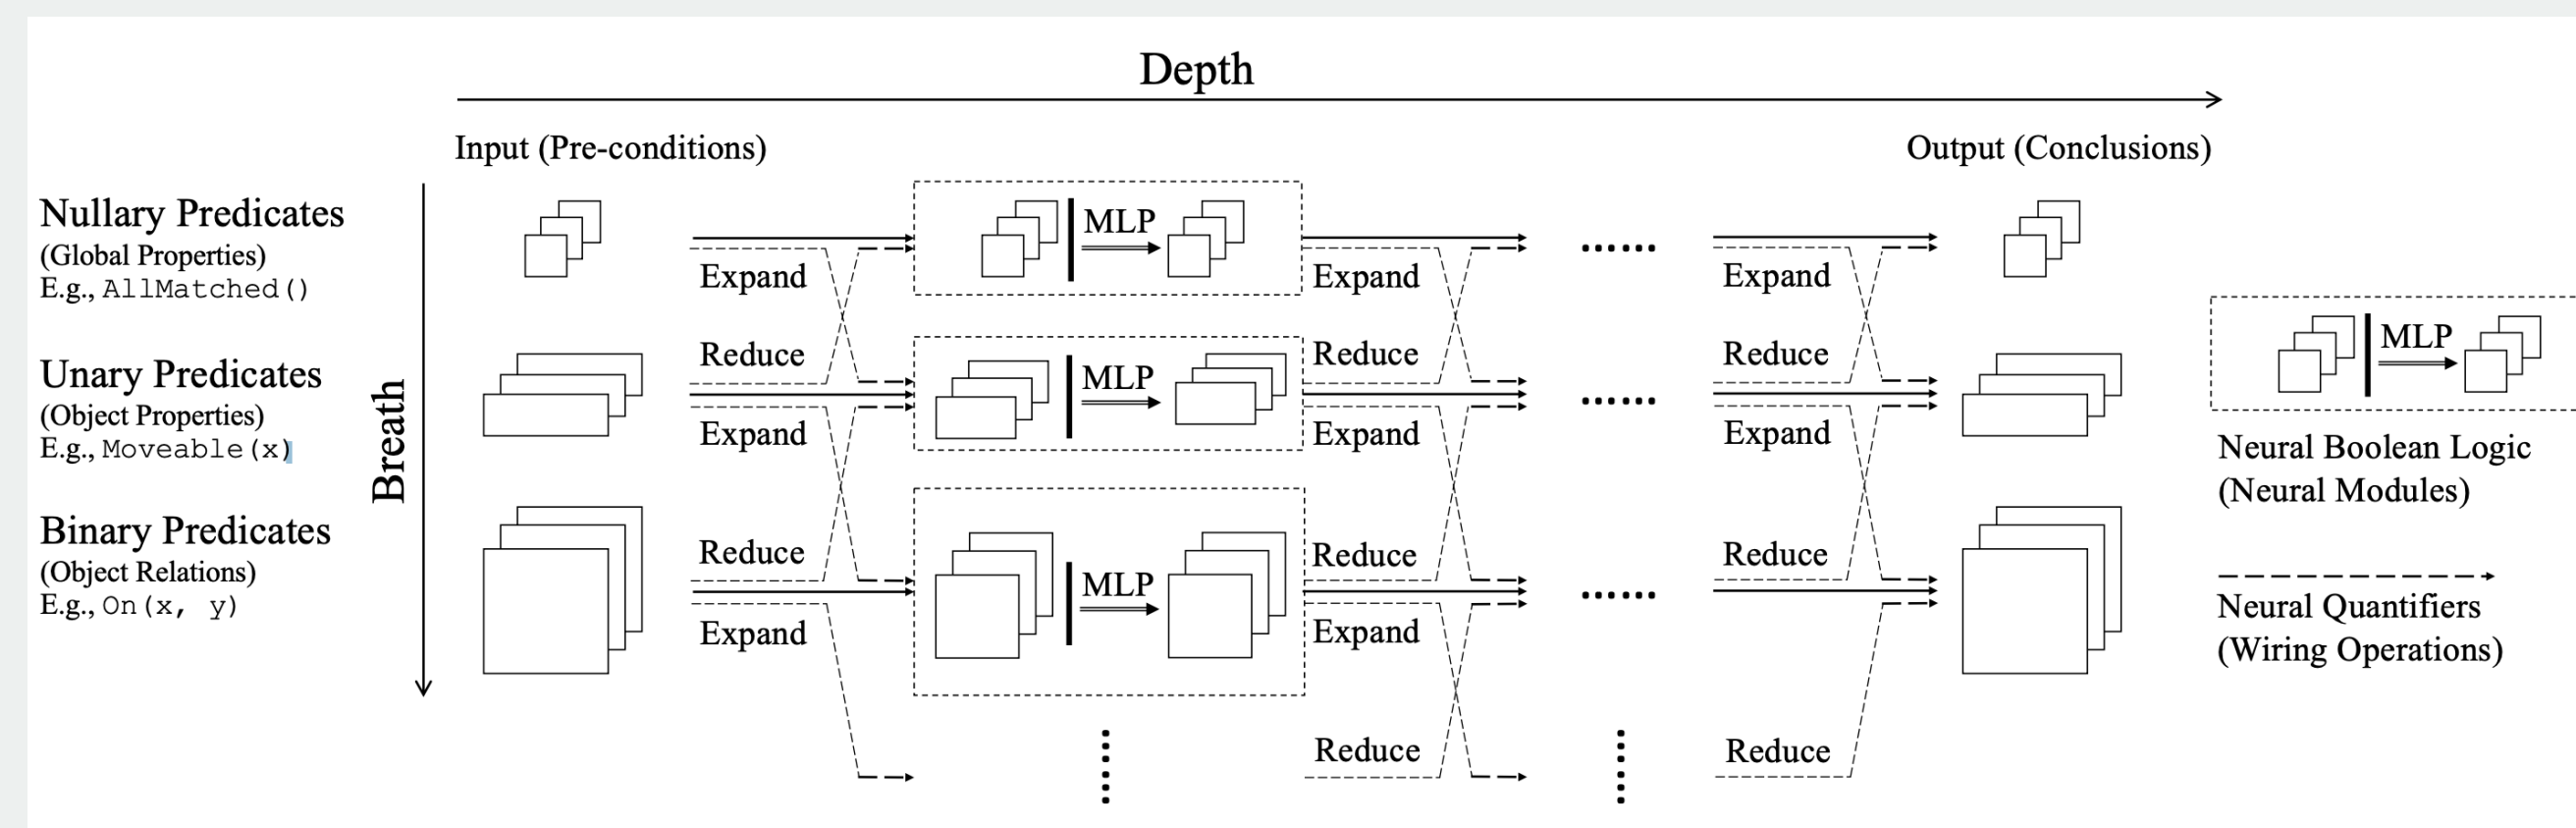
We want to combine RL and domain **independent** heuristics to efficiently learn domain **dependent heuristics**. By learning a specialized heuristic, we hope to plan more efficiently for various problem instances of targeted domain.

## Contributions

1. Generalizing over problem instances by learning **goal** and **problem** conditioned heuristic using neural logic machines [1], and
2. using potential-based reward shaping to efficiently learn a domain **dependent** heuristic as a correction to a domain **independent** heuristic.

## Encoding Domain Dependent Heuristics

We encode the grounded **state** and **goal** predicates with binary N-d arrays. We represent the heuristic with a **neural logic machine** [1].



## Potential-Based Reward Shaping

Using a potential function  $h$ , define a **new reward function**[2]:

$$\hat{r}(s_t, a_t, s_{t+1}) = r(s_t, a_t, s_{t+1}) + h(s_t) - \gamma h(s_{t+1})$$

Theoretically nice approach:

- Optimal policies under the shaped rewards are optimal under the original rewards (and vice versa)

## Evaluation Methodology

- Train on **small** instances (fast and easy), e.g., 2-6 blocks.
- Evaluate on **unseen** and **larger** instances, e.g., 10-50 blocks.
- Evaluate quality of the learned heuristic based on the number of node evaluations.
- Limited to no more than **100,000** node evaluations using **greedy best-first search (GBFS)**.

## References

- [1] Honghua Dong, Jiayuan Mao, Tian Lin, Chong Wang, Lihong Li, and Denny Zhou. “Neural Logic Machines”. In: *ICLR*. 2018.
- [2] Andrew Y Ng, Daishi Harada, and Stuart J Russell. “Policy Invariance Under Reward Transformations: Theory and Application to Reward Shaping”. In: *ICML*. 1999.
- [3] Or Rivlin, Tamir Hazan, and Erez Karpas. “Generalized planning with deep reinforcement learning”. In: *arXiv preprint arXiv:2005.02305* (2020).
- [4] William Shen, Felipe Trevizan, and Sylvie Thiébaux. “Learning Domain-Independent Planning Heuristics with Hypergraph Networks”. In: vol. 30. 2020, pp. 574–584.

## Improving Baseline Heuristics

| domain (total)  | Baselines          |                  |                 | Ours               |                  |                 |
|-----------------|--------------------|------------------|-----------------|--------------------|------------------|-----------------|
|                 | $h^{\text{blind}}$ | $h^{\text{add}}$ | $h^{\text{FF}}$ | $H^{\text{blind}}$ | $H^{\text{add}}$ | $H^{\text{FF}}$ |
| blocks (250)    | 0                  | 5                | 9               | 94                 | 224              | 109             |
| ferry (250)     | 0                  | 0                | 0               | 62                 | 249              | 250             |
| gripper (250)   | 0                  | 0                | 50              | 85                 | 250              | 200             |
| logistics (250) | 0                  | 14               | 77              | 0                  | 108              | 167             |
| miconic (442)   | 17                 | 61               | 0               | 234                | 381              | 442             |
| parking (700)   | 0                  | 105              | 173             | 2                  | 508              | 484             |
| satellite (250) | 0                  | 110              | 63              | 93                 | 115              | 155             |
| visittall (252) | 60                 | 36               | 59              | 192                | 216              | 192             |

Table 2: The number of test instances solved by the learned heuristic but not the baseline and vice versa.

## Acknowledgements

We thank Shen, Trevizan, and Thiebaux and Rivlin, Hazan, and Karpas for kindly providing us their code.