

Motivation and Contributions

Fully Observable Non-Deterministic (FOND) planning models uncertainty through actions with *non-deterministic* effects with uniform probabilities over the actions' effects.

Existing FOND planners are effective and employ a wide range of techniques. However, most of the existing FOND planners are not robust for dealing with both non-determinism and task size.

The main contributions of this work are as follows:

- **Iterative Depth-First Search** (IDFS) algorithms for FOND Planning:
 - Our algorithms are explicitly designed for solving FOND Planning;
 - Use heuristics to make the iterative process effective in FOND;
 - IDFS is effective to deal with different non-deterministic aspects of FOND Planning;

Background and Notation

Definition (FOND Planning Task)

A FOND *planning task* is tuple $\Pi = \langle \mathcal{D}, s_0, s_* \rangle$:

- $\mathcal{D} = \langle \mathcal{F}, \mathcal{A} \rangle$ is a **non-deterministic domain model**, where \mathcal{F} is a set of *fluents*, and a set of **non-deterministic actions** \mathcal{A} . Every $a \in \mathcal{A}$ consists of $a = \langle pre, EFS \rangle$, where:
 - $pre(a)$ represents the **preconditions**; and
 - $EFS(a)$ represents the **set of possible effects** of a ;
- s_0 is the **initial state**;
- s_* is the **goal condition**;

The application of $EFS(a)$ to a state s generates a set of possible successor states $SUCCS(s, a) = \{succ(s, eff) \mid eff \in EFS(a)\}$.

Definition (FOND Planning Solution Policy π)

A *solution* to a FOND planning task Π is a **policy** π , a *partial function* that maps *non-goal states* s into *actions* $a \in \mathcal{A}$.

A *policy* π induces π -trajectories, a non-empty sequence of states $\langle s^0, s^1, \dots, s^k \rangle$, such that $s^{i+1} \in SUCCS(s^i, \pi(s^i))$, $\forall i \in \{0, 1, \dots, k-1\}$.

A *policy* π is *closed* if any π -trajectory starting from s_0 ends either in a *goal state* or in a state defined in the policy π .

Definition (Strong Policy)

A policy π is a *strong policy* for Π if it is closed and no π -trajectory passes through a state more than once.

Definition (Strong Cyclic Policy)

A policy π is a *strong cyclic policy* for Π if it is closed and any π -trajectory starting from s_0 which does not end in a goal state, ends in a state s' such that exists another π -trajectory starting from s' ending in a goal state.




Figure 1. Strong Policy.




Figure 2. Strong Cyclic Policy.

Evaluation Function \mathcal{F} for FOND

^a $g(s)$ represents the search depth from s_0 to s

We define the f -value^a of a state s as:

$$f(s) = g(s) + h(s). \quad (1)$$

We **evaluate the successor states** $SUCCS(s, a)$ using the *Evaluation Function* \mathcal{F}_ξ . \mathcal{F}_ξ uses a function ξ to aggregate the f -values of states in $SUCCS(s, a)$:

$$\mathcal{F}_{\min}(SUCCS(s, a)) \equiv \min_{s' \in SUCCS(s, a)} f(s') \quad (2)$$

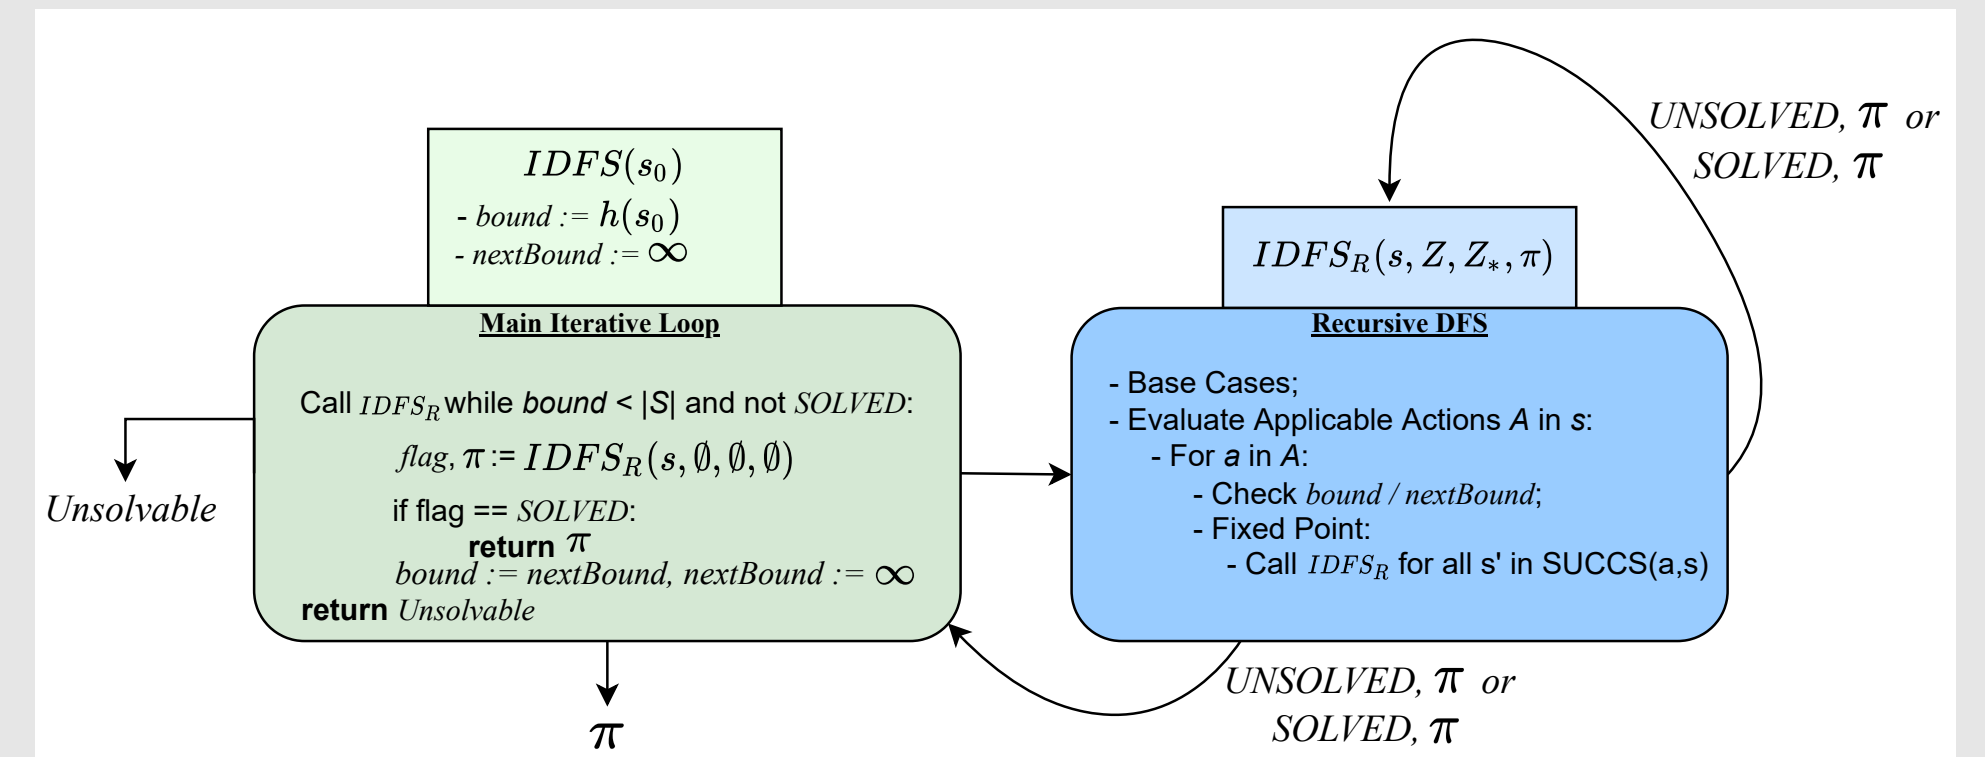
$$\mathcal{F}_{\max}(SUCCS(s, a)) \equiv \max_{s' \in SUCCS(s, a)} f(s') \quad (3)$$

\mathcal{F}_ξ is “pessimistic” when $\xi = \max$, whereas it is “optimistic” when $\xi = \min$.

Iterative DFS for FOND Planning

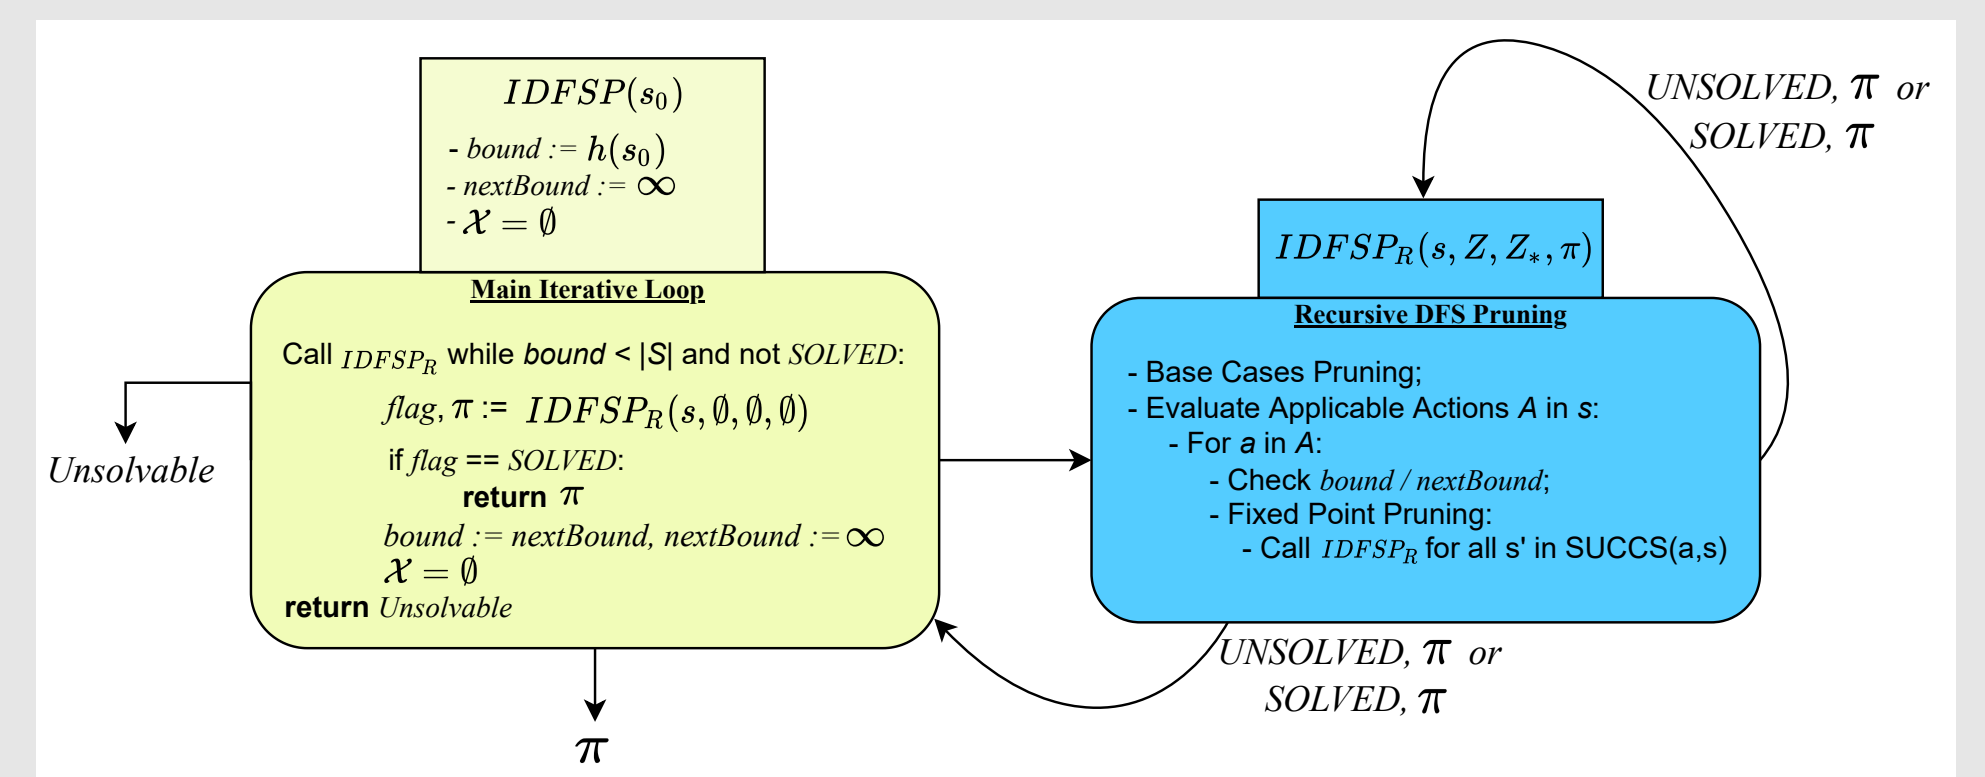
- IDFS performs a series of **bounded** depth-first searches;
- IDFS produces a strong cyclic policy in a **bottom-up way**;
 - It only adds an action to the policy if it determines that the resulting policy with the additional action has the potential to become a strong cyclic policy without exceeding the current search-depth bound.

IDFS



- Z **ancestors** of s ; Z_* **ancestors** of s that has achieved s_* through a π -trajectory;
- **Base Cases:** Check if $s \models s_*$ OR $s \in Z_*$ OR $\pi(s) \neq \perp$. If so, it returns *SOLVED*; And, if $s \in (Z \setminus Z_*)$ (s visited before), if so, it returns *UNSOLVED*.
- **Evaluate Applicable Actions in s :** Check *bound* for $SUCCS(s, a)$.
 - 1 If $\mathcal{F}_\xi(SUCCS(s, a)) > bound$ and $Z_* = \emptyset$: **discards** a on s , and proceeds to the next action.
 - 2 If if $g(n) + 1 > bound$: **discards** a on s , and proceeds to the next action.
- **Fixed Point:**
 - It maps $s \mapsto a$ to π **ONLY** if **ALL** recursive calls on states of $SUCCS(s, a)$ returned *SOLVED*;
 - If not, it **discards** a on s , and proceeds to the next action.

IDFS Pruning (IDFSP)



IDFSP **prunes non-promising** states. E.g., it prunes states whose all successor states for all applicable actions have \mathcal{F} **greater than** the *current bound*.

Experiments and Evaluation

We use two FOND planning benchmarks: IPC-FOND and NEW-FOND.

Our planner is called PALADINUS.

We compare PALADINUS with: MYND, PRP, and FONDSAT.

Planner	Solved Tasks (#590)
PALADINUS IDFS ($\mathcal{F}_{\min}, h^{\max}$)	337
PALADINUS IDFS ($\mathcal{F}_{\min}, h^{\text{FF}}$)	406
PALADINUS IDFS ($\mathcal{F}_{\min}, h^{\text{ADD}}$)	411
PALADINUS IDFS ($\mathcal{F}_{\max}, h^{\max}$)	334
PALADINUS IDFS ($\mathcal{F}_{\max}, h^{\text{FF}}$)	380
PALADINUS IDFS ($\mathcal{F}_{\max}, h^{\text{ADD}}$)	422
FONDSAT	276
PRP (h^{\max})	292
PRP (h^{FF})	412
PRP (h^{ADD})	389
MYND (h^{\max})	180
MYND (h^{FF})	265
MYND (h^{ADD})	289

Table 1. Overall coverage results.

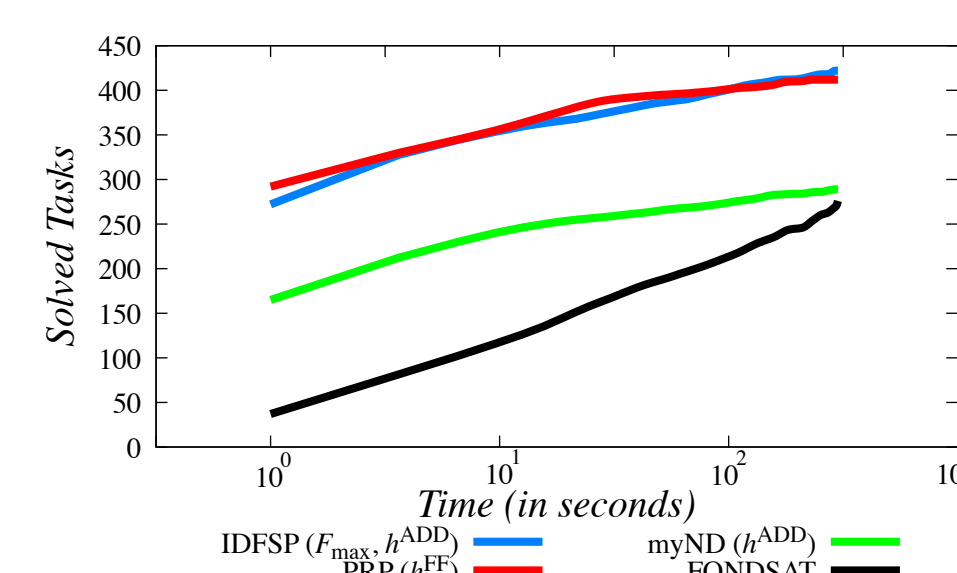


Figure 3. All Benchmarks.




Figure 4. NEW-FOND Benchmarks.